

Deep Learning for Text

Ayoub Bagheri

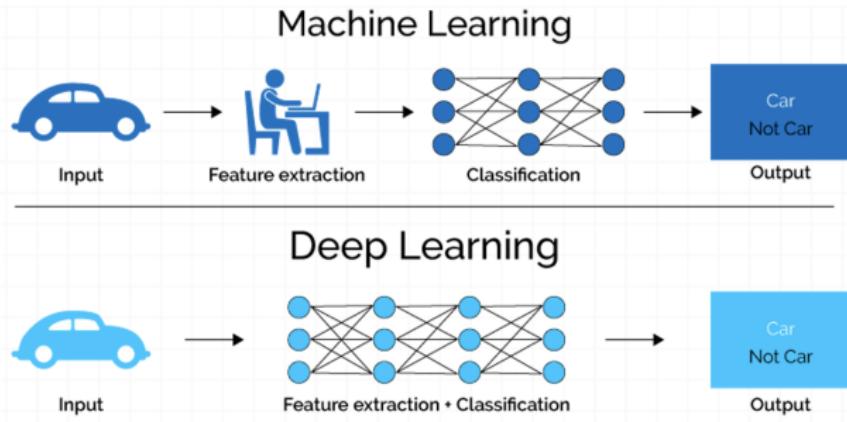
Lecture plan

1. Deep learning
2. Feed-forward neural networks
3. Recurrent neural networks

What is Deep Learning (DL)?

A machine learning subfield of learning representations of data.
Exceptional effective at learning patterns.

Deep learning algorithms attempt to learn (multiple levels of) representation by using a hierarchy of multiple layers.



Deep learning vs neural networks

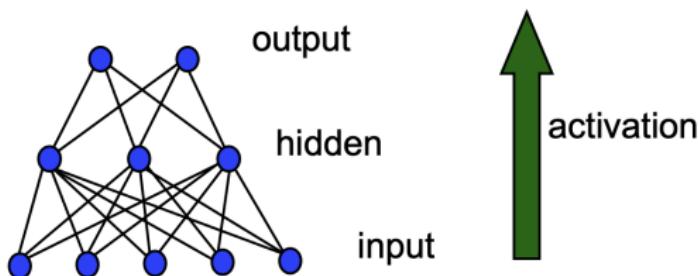
- ▶ Deep learning is only “deep” neural networks, such that with multiple (>2) layers.

Deep learning architectures

- ▶ Feed-forward neural networks
- ▶ Convolutional neural networks
- ▶ Recurrent neural networks
- ▶ Self-organizing maps
- ▶ Autoencoders
- ▶ Transformers: Large Language Models (LLMs)

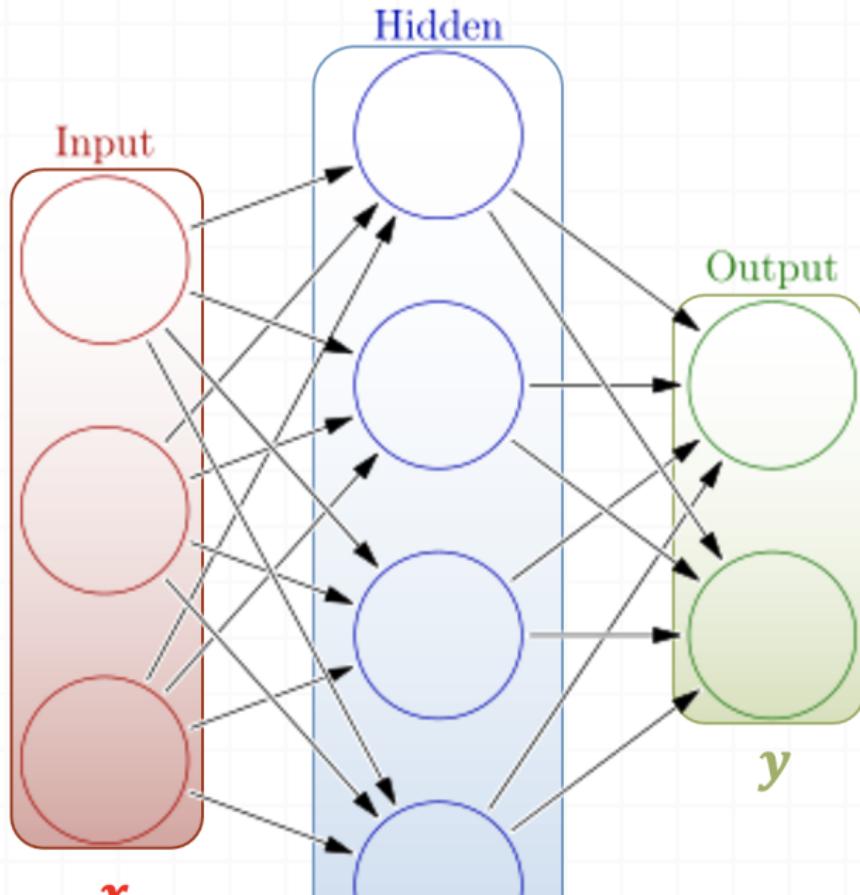
Feed-forward neural networks

- ▶ A typical multi-layer network consists of an input, hidden and output layer, each fully connected to the next, with activation feeding forward.

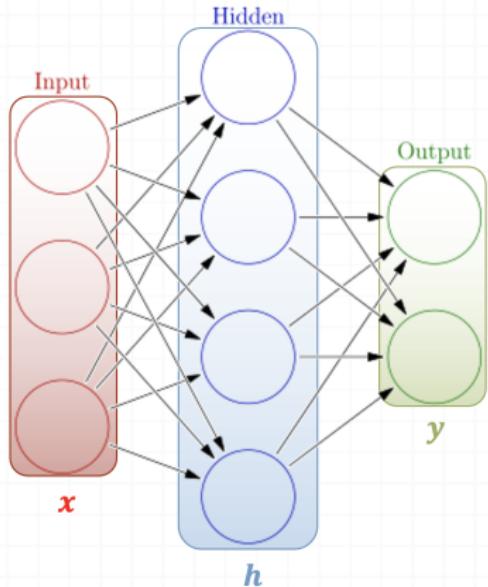


- ▶ The weights determine the function computed.

Feed-forward neural networks



Feed-forward neural networks



Weights

$$h = \sigma(W_1 x + b_1)$$

Activation functions

$$y = \sigma(W_2 h + b_2)$$

$4 + 2 = 6$ neurons (not counting inputs)

$$[3 \times 4] + [4 \times 2] = 20$$
 weights

$$4 + 2 = 6$$
 biases

26 learnable parameters

One forward pass

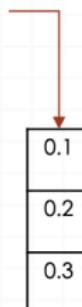
Text (input) representation

TFIDF

Word embeddings

....

0.2	-0.5	0.1
2.0	1.5	1.3
0.5	0.0	0.25
-0.3	2.0	0.0



1.0	3.0	0.025
0.0		



0.95	3.89	0.15
0.37		

very positive
positive
negative
very negative

\mathbf{w}

x_i

b

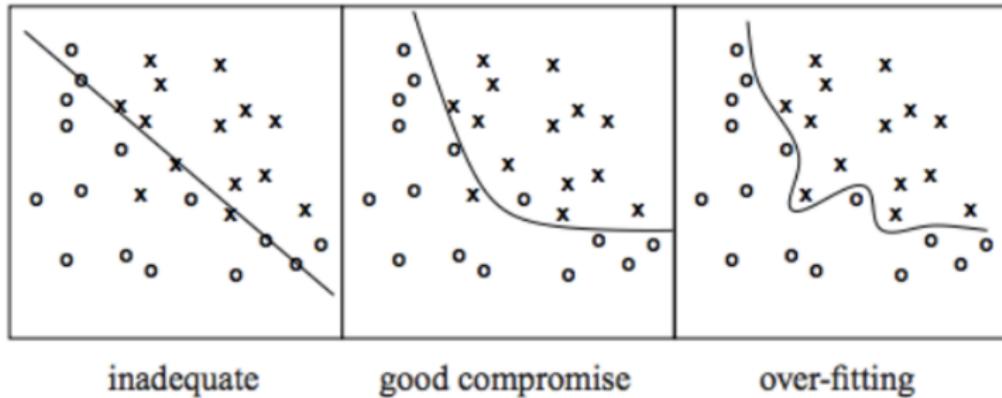
$\sigma(x_i; \mathbf{W}, \mathbf{b})$

Hidden unit representations

- ▶ Trained hidden units can be seen as newly constructed features that make the target concept linearly separable in the transformed space.
- ▶ On many real domains, hidden units can be interpreted as representing meaningful features such as vowel detectors or edge detectors, etc..
- ▶ However, the hidden layer can also become a distributed representation of the input in which each individual unit is not easily interpretable as a meaningful feature.

Overfitting

Learned hypothesis may fit the training data very well, even outliers (noise) but fail to generalize to new examples (test data)

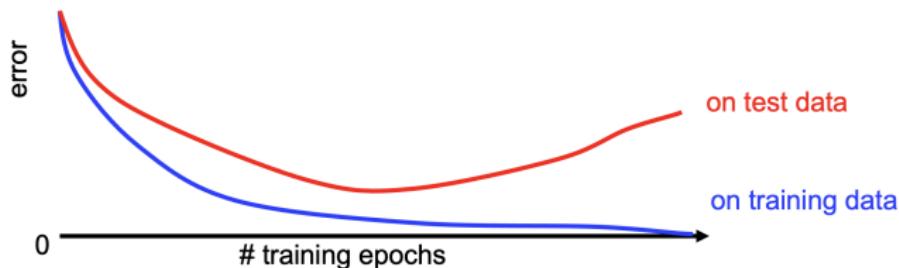


error

<http://wiki.bethanycrane.com/overfitting-of-data>

Overfitting prevention

- ▶ Running too many epochs can result in over-fitting.



- ▶ Keep a hold-out validation set and test accuracy on it after every epoch. Stop training when additional epochs actually increase validation error.
- ▶ To avoid losing training data for validation:
 - ▶ Use internal K-fold CV on the training set to compute the average number of epochs that maximizes generalization accuracy.
 - ▶ Train final network on complete training set for this many epochs.

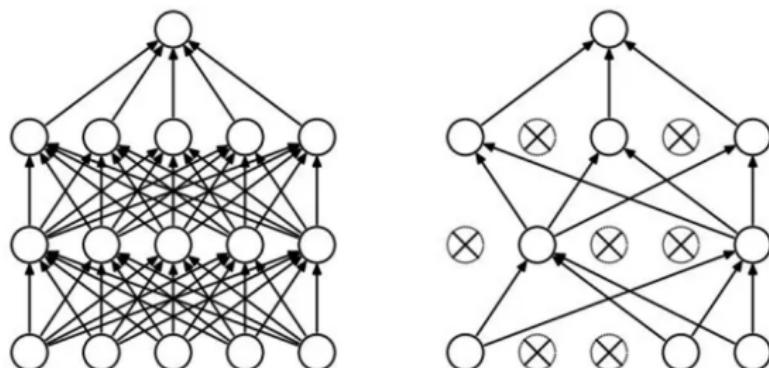
Regularization

Dropout

Randomly drop units (along with their connections) during training

Each unit retained with fixed probability p , independent of other units

Hyper-parameter p to be chosen (tuned)



Srivastava, Nitish, et al. ["Dropout: a simple way to prevent neural networks from overfitting."](#) Journal of machine learning research (2014)

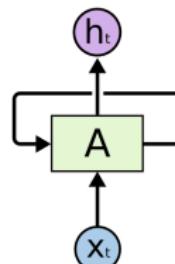
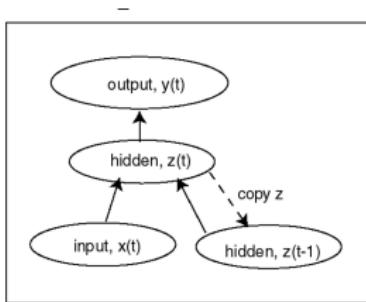
Recurrent Neural Networks

Recurrent Neural Network (RNN)

- ▶ Add feedback loops where some units' current outputs determine some future network inputs.
- ▶ RNNs can model dynamic finite-state machines, beyond the static combinatorial circuits modeled by feed-forward networks.

Simple Recurrent Network (SRN)

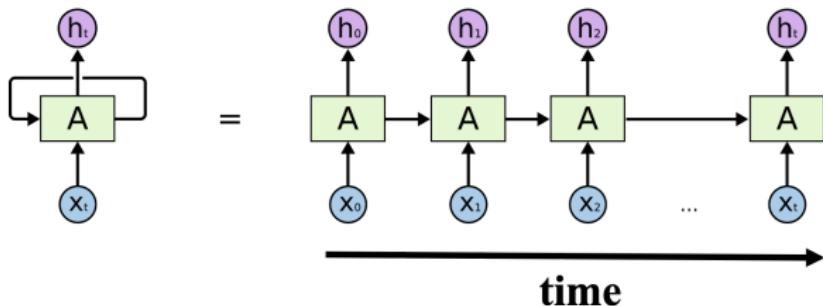
- ▶ Initially developed by Jeff Elman ("*Finding structure in time*," 1990).
- ▶ Additional input to hidden layer is the state of the hidden layer in the previous time step.



<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

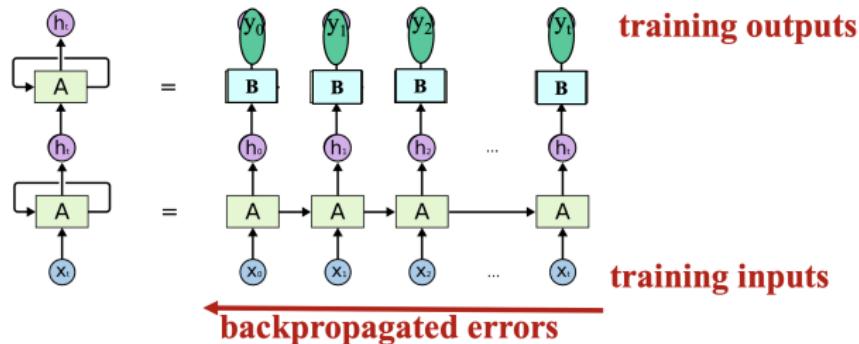
Unrolled RNN

- Behavior of RNN is perhaps best viewed by “unrolling” the network over time.



Training RNNs

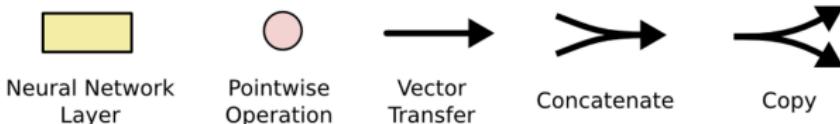
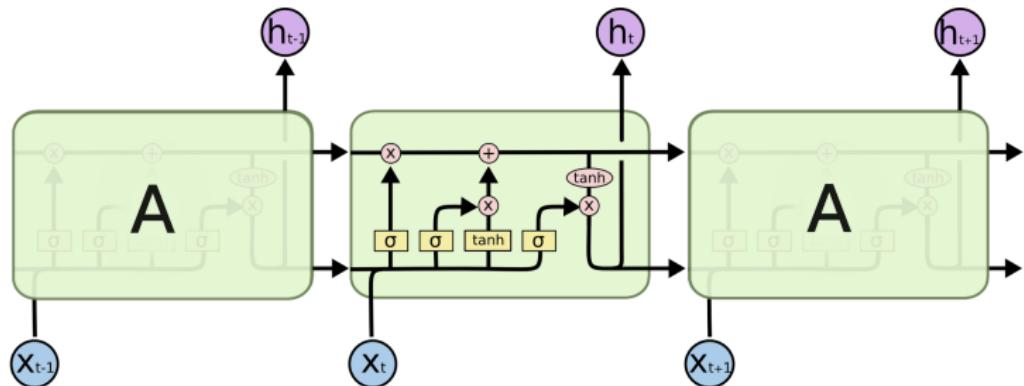
- ▶ RNNs can be trained using “backpropagation through time.”
- ▶ Can be viewed as applying normal backprop to the unrolled network.



Long Short Term Memory (LSTM)

- ▶ LSTM networks, add additional gating units in each memory cell.
 - ▶ Forget gate
 - ▶ Input gate
 - ▶ Output gate
- ▶ Prevents vanishing/exploding gradient problem and allows network to retain state information over longer periods of time.

LSTM network architecture | <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>



In R

```
# Use Keras Functional API
input <- layer_input(shape = list(maxlen), name = "input")

model <- input %>%
  layer_embedding(input_dim = max_words, output_dim = dim_s
                  weights = list(word_embeds), trainable =
  layer_lstm(units = 80, return_sequences = TRUE)

output <- model           %>%
  layer_global_max_pooling_1d() %>%
  layer_dense(units = 1, activation = "sigmoid")

model <- keras_model(input, output)

summary(model)
```

In R

```
## Model: "model"
## _____
##   Layer (type)        Output Shape       Param #  Trainable
## =====
##   input (InputLayer)    [(None, 60)]        0        Y
##   embedding (Embedding) (None, 60, 300)    3000000    N
##   lstm (LSTM)          (None, 60, 80)     121920    Y
##   global_max_pooling1d (GlobalM  (None, 80)        0        Y
##   axPooling1D)
##   dense (Dense)         (None, 1)           81        Y
## =====
## Total params: 3,122,001
## Trainable params: 122,001
## Non-trainable params: 3,000,000
## _____
```

In R

```
# instead of accuracy we can use "AUC" metrics from "tensorflow"
model %>% compile(
  optimizer = "adam",
  loss = "binary_crossentropy",
  metrics = tensorflow::tf$keras$metrics$AUC() # metrics =
)

```

In R

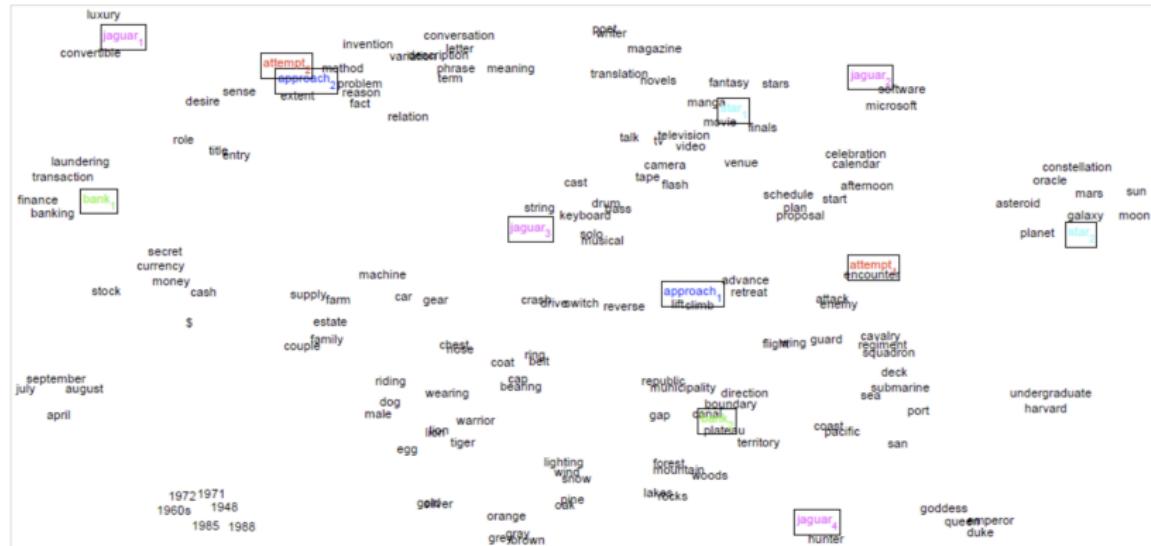
```
history <- model %>% keras::fit(  
  x_train, y_train,  
  epochs = 10,  
  batch_size = 32,  
  validation_split = 0.2  
)
```

Transformers

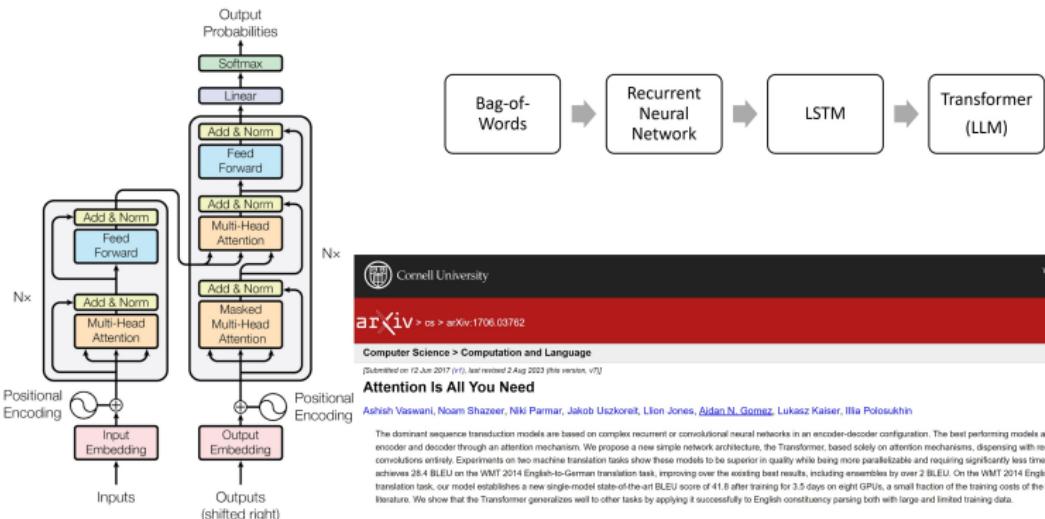
Transformers



Contextual Word Embeddings



Transformers

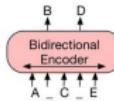


Transformers

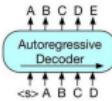
- ▶ A transformer adopts an encoder-decoder architecture.
- ▶ Transformers were developed to solve the problem of sequence transduction, or neural machine translation. That means any task that transforms an input sequence to an output sequence.
- ▶ More details on the architecture and implementation:
 - ▶ <https://arxiv.org/abs/1810.04805>
 - ▶ <http://nlp.seas.harvard.edu/2018/04/03/attention.html>
 - ▶ <https://jalammar.github.io/illustrated-transformer/>

Transformer foundation models: BERT, GPT, BART

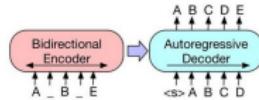
- **BERT: Bidirectional Encoder Representations from Transformers.**
 - *Masked word prediction, text representation*
- **GPT: Generative Pre-trained Transformer.**
 - *Next word prediction, text generation, chat*
- **BART = “BERT+GPT”:** Bidirectional encoder and Auto-Regressive decoder Transformers.
 - *Noised text reconstruction, summarization, translation, spelling correction*



(a) BERT: Random tokens are replaced with masks, and the document is encoded bidirectionally. Missing tokens are predicted independently, so BERT cannot easily be used for generation.



(b) GPT: Tokens are predicted auto-regressively, meaning GPT can be used for generation. However words can only condition on leftward context, so it cannot learn bidirectional interactions.



(c) BART: Inputs to the encoder need not be aligned with decoder outputs, allowing arbitrary noise transformations. Here, a document has been corrupted by replacing spans of text with a mask symbol. The corrupted document (left) is encoded with a bidirectional model, and then the likelihood of the original document (right) is calculated with an autoregressive decoder. For fine-tuning, an uncorrupted document is input to both the encoder and decoder, and we use representations from the final hidden state of the decoder.

BERT: Bidirectional Encoder Representations from Tranformers

BERT: Bidirectional Encoder Representations from Tranformers

Transformers

- ▶ ChatGPT: <https://chat.openai.com/>
- ▶ Write with Transformer: <https://transformer.huggingface.co/>
- ▶ Talk to Transformer: <https://app.inferkit.com/demo>
- ▶ Transformer model for language understanding:
<https://www.tensorflow.org/text/tutorials/transformer>
- ▶ Pre-trained models:
https://huggingface.co/transformers/pretrained_models.html

ChatGPT (5-min exercise)

- ▶ Go to <https://chat.openai.com/> and login
- ▶ How many parameters has chatgpt-3 model been trained on?
- ▶ How many parameters has chatgpt-4 model been trained on?
- ▶ What is the next generation NLP?
- ▶ Suppose we want to build an application to help a user buy a car from textual catalogues. The user looks for any car cheaper than \$10,000.00. Assume we are using the following data: `txt <- c("Price of Tesla S is $8599.99.", "Audi Q4 is $7000.", "BMW X5 costs $900")`. Could you give me a regular expression to do this in R?

Summary

Summary

- ▶ Deep learning
- ▶ Feed-forward neural networks
- ▶ Recurrent neural networks
- ▶ State-of-the-art LLMs

Practical 8